

Achieving Safe Control Online through Integration of Harmonic Control Lyapunov-Barrier Functions with Unsafe Object-Centric Action Policies

Marlow Fawn

Matthias Scheutz

Human-Robot Interaction Laboratory
Tufts University
Medford MA 02155, USA

marlow.fawn@tufts.edu

matthias.scheutz@tufts.edu

We propose a method for combining Harmonic Control Lyapunov–Barrier Functions (HCLBFs) derived from Signal Temporal Logic (STL) specifications with any given robot policy to turn an unsafe policy into a safe one with formal guarantees. The two components are combined via HCLBF-derived safety certificates, thus producing commands that preserve both safety and task-driven behavior. We demonstrate with a simple proof-of-concept implementation for an object-centric force-based policy trained through reinforcement learning for a movement task of a stationary robot arm that is able to avoid colliding with obstacles on a table top after combining the policy with the safety constraints. The proposed method can be generalized to more complex specifications and dynamic task settings.

1 Introduction

Open-world environments pose many challenges for autonomous robots as unexpected events or task modulations can make learned robot behavior inapplicable or obsolete. Consider, for example, a robot that has learned to autonomously perform a sorting task on a table top without any human interventions when a human co-worker steps in to help with finishing the task. This change in task environment now requires the robot to avoid colliding with the human whose arms are extended into the robot’s work space and are dynamically changing position. Even if the robot has the perceptual capability to detect and track the human’s arms and hands, its trained action policy does not provide a way to account for the motion constraints they impose. Or consider a delivery robot in a warehouse that has an optimized policy for traversing indoor spaces when dynamic constraints are imposed on where it can drive (e.g., because parts of the floor are painted). In both cases it would be important for the robot to be able to adjust its behavior on the fly to account for the changes in the task environment, ideally with formal guarantees that these changes are *safe*.

State-of-the-art learning methods for training robots on tasks such as Reinforcement Learning (RL) enable robots to learn complex behaviors through trial and error, in some cases with safety guarantees such as in the case of “safe RL” [4], but the trained behaviors cannot be easily modified or adapted to new circumstances or safety constraints. In this paper, we utilize analytical methods, such as Control Lyapunov Functions (CLFs) and Control Barrier Functions (CBFs), which offer formal guarantees of stability and safety, to enable the adaptation of previously learned control policies to changes in the task environment. Specifically, we propose to use *Signal Temporal Logic* (STL) to describe the changed task constraints from which we then derive Harmonic Control Lyapunov–Barrier Functions (HCLBFs) whose gradients can drive the system toward its goals while avoiding unsafe regions. We then integrate the safety constraints with the robot’s control policy through a shared velocity representation that allows

the HCLBF field to modulate the learned policy without retraining. Importantly, due to the way the components are integrated, the resultant behavior is provably safe.

There exists a small body of work integrating CBFs with STL. In particular, [12, 11] both provide robust STL CBF integration for agent navigation. Our primary differences are twofold: (1) We expand on these method by integrating Lyapunov-like functionality, further expanding the valid STL definitions and (2) We use these CLBFs for *synthesis* with CLBF control laws, allowing for greater task flexibility through RL. With our method, we ultimately hope to integrate control laws with temporal logic, encompassing these two methods in a larger framework that allows for dynamically safe control of pre-existing behaviors.

The rest of the paper is structured as follows: After briefly reviewing some formal preliminaries, we introduce the proposed method for deriving STL-based HCLBFs using a restricted version of STL and Laplace solver for generating HCLBFs on a grid rather than a continuous spatial environment. Next, we introduce the demonstration environment in which we train an “object-centric force-based RL policy” (e.g., [3]), i.e., a policy that represents forces applied to object in the direction of its intended motion, in a simple proof-of-concept arm movement task in the 2D plane from a start to a goal position. We then show how we can integrate both components into a unified control loop with HCLBF certificates so that the robot avoids colliding with unexpected obstacles. Finally, we briefly discuss the advantages and limitations of our proposed approach as well as directions for future work, and conclude with a brief summary of what we accomplished.

2 Preliminaries

We first briefly review three different formalisms: Signal Temporal Logic (STL), Harmonic Control Lyapunov–Barrier Functions (HCLBFs), and Soft Actor Critic (SAC) Reinforcement Learning (RL).

2.1 Signal Temporal Logic (STL)

Signal Temporal Logic (STL) extends linear temporal logic (LTL) which is frequently used as a specification language for robot behavior (e.g., [7]) by allowing predicates to range over real values (e.g., spatial positions) in real-time. I.e., we assume we are given signals of the form $x_1(t), \dots, x_n(t)$ which we can use in atomic positions p of the form

$$p := f(x_1(t), \dots, x_n(t)) \geq 0$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is a *constraint* (e.g., see [2]). Atomic positions can be modulated with two temporal operators $\mathbf{G}_{[a,b]}$ and $\mathbf{F}_{[a,b]}$ inherited from Metric Temporal Logic (MTL), meaning “for all times in $[a,b]$ ” and “there is a time in $[a,b]$ ”, respectively, thus allowing for the specification of both temporal and spatial constraints.

The satisfaction of an STL formula φ by a signal $x = (x_1, \dots, x_n)$ at time t is then defined as follows:

- $(x, t) \models p \leftrightarrow f(x_1(t), \dots, x_n(t)) \geq 0$
- $(x, t) \models \neg \varphi \leftrightarrow (x, t) \not\models \varphi$
- $(x, t) \models \varphi \wedge \psi \leftrightarrow (x, t) \models \varphi \wedge (x, t) \models \psi$
- $(x, t) \models \varphi \vee \psi \leftrightarrow (x, t) \models \varphi \vee (x, t) \models \psi$
- $(x, t) \models \mathbf{F}_{[a,b]} \varphi \leftrightarrow \exists t' \in [t + a, t + b] : (x, t') \models \varphi$

For our purposes here, we will focus on a small subset of STL, namely the set of atomic formulas in the scope of one of the temporal operators and their conjunctions (e.g., $\mathbf{G}_{[a,b]}f(x_1(t)) > 0 \wedge \mathbf{F}_{[a',b']}f'(x_1(t), x_2(t)) > 0$).

2.2 Harmonic Control Lyapunov–Barrier Functions (HCLBFs)

Control Lyapunov Barrier Functions (CLBFs) define a framework for combining the reachability of *Control Lyapunov Functions* with the forward-invariance of *Control Barrier Functions* (e.g., [10]). Given the control function:

$$\dot{s}(t) = f(s(t), u(t)), \quad s \in \mathbb{R}^n, u \in \mathbb{R}^m,$$

a CLBF $V : \mathcal{S} \rightarrow \mathbb{R}$ is defined as the constraints on that function such that: (i) V decreases along trajectories, which drives the system towards a goal and (ii) V has high magnitude in unsafe regions, preventing constraint violation.

We will use the definition for HCLBFs provided by Mukherjee et al. [9]. Let $\mathcal{S} \subset \mathbb{R}^n$ denote the safe set, $\mathcal{S}_{\text{goal}} \subset \mathcal{S}$ the goal region, and $\mathcal{S}_{\text{unsafe}} \subset \mathbb{R}^n \setminus \mathcal{S}$ the unsafe set. A function $V \in C^2(\mathcal{S}) \cup C^1(\partial\mathcal{S})$ is a *harmonic CLBF* if it satisfies:

1. $\nabla^2 V(s) = 0 \quad \forall s \in \mathcal{S}_{\text{safe}};$
2. $V(s) = 0 \quad \forall s \in \mathcal{S}_{\text{goal}};$
3. $V(s) = c \quad \forall s \in \partial\mathcal{S} \cup \mathcal{S}_{\text{unsafe}},$

for some constant $c > 0$. Intuitively,

1. Ensure that the field is never “flat” in safe regions, meaning the system will never stall
2. Ensure the goal is always the “lowest” point in the state space, meaning the system will be drawn to it
3. Ensure the obstacles are always the “highest” points in the state space, meaning the system will be repelled from them

Formally, V is the solution to the boundary value problem defined by Laplace’s equation on \mathcal{S} with Dirichlet boundary conditions set by the goal and unsafe sets. In other words, we are solving for some function V over the state space \mathcal{S} which contains our goal, safe, and unsafe regions, such that the position of an object following V as a trajectory will never overlap with unsafe regions, and will ultimately be driven to goal regions.

We will, in our proof-of-concept implementation in Section 4, approximate this continuous PDE on a finite grid for tractability.

2.3 Reinforcement Learning

Robot tasks and thus task-based control problems are commonly modeled as either fully or partially observable Markov Decision Process (MDP) $(\mathcal{S}, \mathcal{A}, P, r, \gamma)$ where \mathcal{S} is the state space, \mathcal{A} the action space, P the transition function, $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ the reward function, and $\gamma \in (0, 1]$ the discount factor. A stochastic policy $\pi_\theta(a | s)$ with parameters θ induces a distribution over trajectories, and is optimized to maximize the expected discounted return

$$J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right].$$

A common method for learning a policy π_θ is Q-learning which learns a Q-function $Q(s, a)$ for every state $s \in \mathcal{S}$ and action $a \in \mathcal{A}$ through (possibly unsafe) explorations of the state space where $Q(s, a)$ is the expected discounted sum of future rewards received by performing action a in state s . The Q-function is updated as follows:

$$Q \leftarrow Q + \alpha \left[\left(\mathbb{R} + \gamma \max_{a' \in \mathcal{A}} Q(a', s') - Q(a, s) \right) \right]$$

where $\alpha \in (0, 1]$ is the learning rate. Q-learning uses an ε -greedy policy for exploring new states in which the agent selects a random action with probability ε rather than performing the action prescribed by the largest q -value.

The ‘‘Soft Actor Critic’’ (SAC) version of RL, consisting of an ‘‘actor’’ and a ‘‘critic’’ (often more than one critic), is an ‘‘off-policy’’ method that integrates ‘‘maximum entropy’’ reinforcement learning to maximize both the expected reward as well as the entropy of the policy. The actor adds an entropy term at state s to the expected discounted return:

$$\beta \cdot -\mathbb{E}_a \sim \pi_\theta \log \pi(a|s)$$

where β scales the entropy term. The critic, in turn, computes and tracks $Q(a, s)$.

3 STL-Based HCLBFs

We start by defining the permissible formal constraints in STL, followed by a description of our HCLBF implementation. We then explain how we use STL to generate HCLBF constraints, with a brief mention of the particular restrictions for our case based on the resulting potential field.

3.1 STL Definitions

Previous work has shown that CLBFs can be used to satisfy STL formulae for a dynamic system (e.g., [13, 8]). We use a simplified version of this work to construct boundaries for HCLBFs using a small fragment of STL which is defined as follows:

- **Atomic formulas p :** We consider atomic formulas p of the form

$$x \bowtie c \quad \text{or} \quad y \bowtie c,$$

where $\bowtie \in \{\geq, >, =\}$, c is a constant, and x and y are components of a spatial coordinate $\langle x, y \rangle$. All atomic formulas are allowable formulas.

- **Negation:** Only atomic formulas p can be negated, $\neg p$. Negated and non-negated atomic formulas are considered ‘‘allowable formulas’’.
- **Boolean composition:** Allowable formulas φ and ψ can be combined through (\wedge) , $\varphi \wedge \psi$; no other combination such as disjunction (\vee) or implication (\Rightarrow) is allowed.
- **Temporal operators:** Allowable formulas can get temporally quantified with a single bounded ‘‘always’’ or ‘‘eventually’’ operator:

$$\mathbf{G}_{[t_1, t_2]} \psi \quad \text{or} \quad \mathbf{F}_{[t_1, t_2]} \psi,$$

where $0 \leq t_1 < t_2$; nesting of temporal operators is not allowed.

Examples of formulas in the restricted STL fragment are $F_{[t_1, t_2]}(x > 0 \wedge \neg x > 2)$ or $G_{[t_1, t_2]}(x > 0 \wedge y > 0)$. The fragment allows for specifying intervals during which spatial constraints need to hold, e.g., in the example from the introduction one such specification would be to impose that the robot not move into a spatially defined bounding box around the human’s arms *while* the human is operating in the robot’s space, which may need to be updated dynamically. The specification also allows for specifying constraints that will occur within a particular interval (e.g., that there will be a moment when operating in a particular region is forbidden). Note that we interpret constants here as making spatial references to discrete grid coordinates in the robot’s working space, but nothing critical hinges on it (we will later extend this to continuous domains through interpolation outside STL).

3.2 HCLBF Definitions

To reiterate from our section 2.2, in theory, HCLBFs are defined as scalar functions that solve a boundary value problem on a continuous safe set. Formally, we define

$$V : \mathbb{R}^2 \rightarrow [0, 1],$$

where the associated velocity field is obtained by the negative gradient

$$U_\Phi(x, y) = -\nabla V(x, y),$$

which provides directions of motion that flow toward goals while avoiding unsafe regions. Intuitively, an HCLBF is a single function that defines a stable policy for how to move across a space such that certain constraints will be followed with formal guarantees.

In our implementation, we discretize the HCLBF onto a finite grid for tractability and to align with our discrete STL construction of spatial restrictions. The potential field V is thus stored at integer grid points $(i, j) \in \mathbb{Z}^2$, and gradients are approximated using finite differences.

Similar to our STL restrictions, we impose some limitations on the form of the HCLBF that can be relaxed in future work:

- In formal HCLBFs, goal and unsafe regions cannot overlap. However, in real-world scenarios this happens frequently (e.g. an obstacle is placed or moves over a known goal region). As our goal is to prioritize safety, we allow the unsafe regions to overrule goal regions. While this leads to unreachable scenarios, it ensures that all scenarios are fully safe.
- We ignore settings where the agent is already in collision with an obstacle. This is solvable with heuristics driving the agent away from the obstacle region, but is irrelevant for our intended use of fully safe control.
- At least one goal cell is required, otherwise the velocity field degenerates to trivial flow around unsafe regions.
- The current formulation is grid-based, and gradients are approximated numerically as continuous extensions would require analytic PDE solutions or finer interpolation, which is computationally expensive.

In sum, the use of HCLBF with STL is best understood as a potential field encoding of the STL specification: It takes in a spatial state (x, y) and outputs a velocity vector $U_\Phi(x, y)$ that balances progress toward the goal against avoidance of unsafe regions (see Fig. 1 for examples).

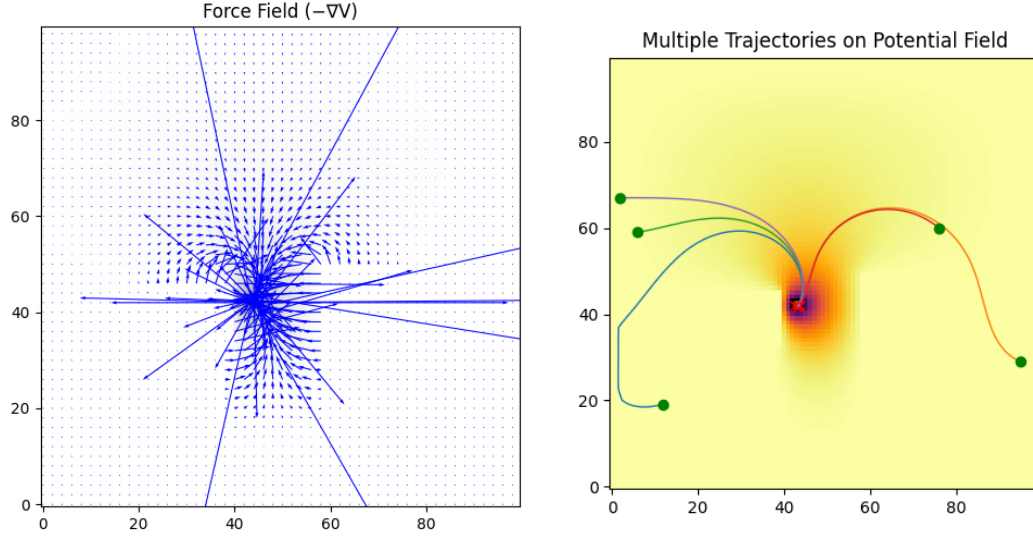


Figure 1: HCLBF field and gradient (left) and sample trajectories demonstrating safe goal reaching (right).

3.3 Laplace Solution for HCLBFs

To generate an HCLBF from a given STL formula, we translate the logical specification into spatial boundary conditions and then solve a Laplace equation on the workspace grid. The STL construction produces two sets, a goal region G and an unsafe region U , which can be converted into Dirichlet boundary conditions for the scalar potential V :

$$V(x,y) = 0 \quad \text{for } (x,y) \in G, \quad V(x,y) = 1 \quad \text{for } (x,y) \in U.$$

On the remaining safe cells, we enforce harmonicity:

$$\nabla^2 V(x,y) = 0,$$

which yields a unique smooth interpolation between goal and unsafe regions. In practice, V is computed on the discrete grid using an iterative relaxation method until convergence.

4 Proof-of-concept Setting

To demonstrate how the above method works to safeguard behavior of autonomous robots when there are changes to a task environment, we consider a simple robot arm navigation task. In this task, the robot arm has to traverse the surface of table top without any obstacles moving from an initial position to a goal position. We start by training an RL policy, as would be commonly done for robot tasks (usually more complex ones, of course) using an object-centric force-based policy (e.g., [3]).

4.1 Policy Learning

We train the object-centric policy with a Soft Actor–Critic (SAC) RL algorithm on a rigid “virtual cube”, i.e., a proxy for the motion of the robot, learning how to move towards randomly sampled goals in task

space. Since training occurs in *object mode*, i.e., forces applied directly to the cube, we can keep the policy independent of the particular robot and robot morphology and thus allow for the trained policy to operate on any robot for which an inverse kinematics model exists (e.g., see [3] for details on the method and a discussion). We will then later use the same learned representations when deploying the policy through an “Operational Space Control” (OSC) [6] interface.

Note that the policy observes the object’s Cartesian position and velocity,

$$s_t = [x(t), v(t)] \in \mathbb{R}^{2n},$$

and outputs a task-space force $a_t = F_t \in \mathbb{R}^n$ (in our experiments $n = 2$ for planar motion). We use a dense shaping signal

$$r_t = -\|x(t) - x_{\text{goal}}\|_2 - 0.01,$$

with a success bonus and out-of-bounds penalty.

In sum, using force as the action space achieves three goals: (i) it keeps the policy object-centric and robot-agnostic, (ii) it produces commands in a common kinematic space that can be directly used with HCLBFs, and (iii) it leverages existing OSC infrastructure for low-level actuation without requiring direct torque control.

4.2 Robot Controller

We assume a rigid, no-slip grasp between the robot’s end effector (EE) and the controlled object, so that the object twist equals the EE twist, $V_{\text{obj}} = V_{\text{ee}}$. This allows forces applied to the object in training to be mapped directly and consistently to EE commands on the robot.

We convert the force F_d outputs of learned policy (on the cube) to a desired EE velocity $U_\pi(x)$ using an admittance model $A(x)$ (e.g., [5]), so that commanded forces translate to proportional velocities in task space. For our implementation, the simplified admittance model where the force F_d is scaled by a constant gain to produce a velocity $U_\pi(x) = \alpha F_d$ is sufficient but nothing hinges on it and more complex models are possible. This model preserves the direction encoded by the policy while providing a tuning for integration into the OSC. It also allows for future substitutions with more complex or kinematically accurate models without affecting the safety of the system.

Since the robot’s OSC accepts Cartesian pose targets, we integrate the commanded velocity over the control period Δt to update the EE position:

$$p_{t+1} = p_t + U_\pi(x_t) \Delta t,$$

with orientation held fixed.

5 Control Synthesis

We now demonstrate how we can combine the HCLBF safety constraints with the learned RL policy. And note that while we are using RL here, as it is commonly used in robotics, in principle any OSC controller could be used in its place (e.g., online motion planners); These controllers would maintain the safety guarantees provided by the HCLBF filtering regardless of underlying controller.

Let

$$u = U_\pi(x) := \mathcal{A}(x) F_\pi(x)$$

denote the nominal velocity command obtained from the policy via the admittance map $\mathcal{A}(x)$. The HCLBF scalar function $V(x)$ and gradient $\nabla V(x)$ are evaluated by mapping the current world position to grid indices and interpolating both the potential and its finite-difference derivatives, converted to world coordinates.

The safety condition is expressed as the barrier inequality

$$\nabla V(x)^\top \tilde{u} \leq -k_\alpha V(x),$$

where $k_\alpha > 0$ is a constant gain. If the nominal u satisfies this inequality, it is accepted. Otherwise, we solve the projection problem

$$u^* = \arg \min_{\tilde{u} \in \mathbb{R}^2} \|\tilde{u} - u\|_2^2 \quad \text{s.t.} \quad \nabla V(x)^\top \tilde{u} \leq -k_\alpha V(x).$$

This ensures the final action is the closest possible safe modification of the policy output.

The controller thus applies the HCLBF safety filter to minimally correct the learned policy's action as follows:

1. Query the learned policy to obtain a task-space force $F_\pi(x)$.
2. Map the force to a task-space velocity $u = U_\pi(x)$ using an admittance model.
3. Interpolate the HCLBF scalar field $V(x)$ and its gradient $\nabla V(x)$ at the current end-effector position.
4. Check the barrier inequality $\nabla V(x)^\top u \leq -k_\alpha V(x)$. If satisfied, keep u . If violated, project u onto the safe half-space to obtain the corrected velocity u^* .
5. Package the final command with vertical and gripper components, integrate over the control period, and send the pose target to the robot's OSC controller.

Note that the above algorithm ensures that safety constraints are enforced at every step without retraining or modifying the learned policy.

We then extend corrected velocity u^* with vertical and gripper components to form

$$u_{\text{cmd}} = [u_x^*, u_y^*, z_{\text{hold}} - z_{\text{ee}}, g_{\text{default}}],$$

which is integrated over the control period Δt to update the end-effector target:

$$p_{t+1} = p_t + u^* \Delta t.$$

The new position p_{t+1} is passed to the OSC_POSITION controller, which computes joint commands. The environment (simulation or hardware) then advances one step, and the cycle repeats until termination. Fig. 2 shows the experimental evaluation setting in the *RoboSuite* robotics simulation environment.

6 Discussion

The proof-of-concept demonstration shows the viability of the proposed method for generating safety guardrails for unsafe policies online. However, the online computation speed will depend on multiple factors, including the complexity of the safety specifications, both in terms of the spatial regions but also the temporal complexity and duration of the intervals as well as the granularity of the discretization and convergence tolerance of the solver. For constraints with very short intervals arriving in real-time, the time to calculate the safety velocity field might exceed the constraint time. Also, extensions to




Figure 2: An image from the implementation of the above method in the *RoboSuite* simulation environment showing the task arena with obstacles (red) and the robot arm’s goal point (green).

higher-dimensional workspaces will require efficient solvers which can generate sufficient approximations quickly. On the other hand, if the setting is static or does not change frequently, there will be sufficient time to calculate the potentials and the computational effort and time to do so will likely be significantly shorter than adapting the robot’s control through RL or training it from scratch.

It is also interesting to consider the proposed approach in light of work in safe RL (e.g., [4]) where a “safety shield” is provided initially that bounds the region within which the robot can explore and learn its policy. The current approach could be used in a similar way during regular RL training when the robot needs to learn an action policy as the “exploration” part of RL might make the robot attempt actions that lead to unsafe regions – using the safety HCLBF the robot would be prevented from carrying out unsafe actions.

The current method is not without several limitations that will be addressed in future work. For one, we used a very limited subset of STL for specifying only spatial constraints in order to simplify the construction of HCLBFs. Extending the usable subset of STL to richer logical structures, including disjunctions and nested temporal operators, would allow for more complex specifications of time-varying goals and constraints. It would also be important to explore the use of continuous PDE solvers to be able to replace the current grid-based Laplace approximation, which would enable higher-dimensional workspaces and real-time updates. Existing approaches to explore include sum-of-squares methods [1] and neural CBFs [14]. Finally, the current implementation follows theory to create fully safe navigation, and initial trials have proven successful. However, this must be evaluated formally with quantitative analyses to demonstrate safety in regular operation.

7 Conclusion

In this paper we proposed a new way of adapting an existing robot policy (e.g., trained through reinforcement learning) to meet new (potentially dynamic) motion constraints without requiring retraining of the policy. The proposed method utilizes Harmonic Control Lyapunov–Barrier Functions generated from formal constraints in Signal Temporal Logic which are then combined with the learned policy in a shared velocity space. The method enables online adaptations of robot behavior that inherit the formal guarantees of Harmonic Control Lyapunov–Barrier Functions, providing a unified control framework that preserves both task performance and safety.

References

- [1] Hongkai Dai, Chuanrui Jiang, Hongchao Zhang & Andrew Clark (2024): *Verification and Synthesis of Compatible Control Lyapunov and Control Barrier Functions*, doi:10.48550/arXiv.2406.18914. Available at <https://arxiv.org/abs/2406.18914>.
- [2] Alexandre Donzé (2013): *On Signal Temporal Logic*. In Axel Legay & Saddek Bensalem, editors: *Runtime Verification*, Springer Berlin Heidelberg, Berlin, Heidelberg, pp. 382–383, doi:10.1007/978-3-642-40787-1_27.
- [3] Shijie Fang, Wenchang Gao, Shivam Goel, Christopher Thierauf, Matthias Scheutz & Jivko Sinapov (2025): *FLEX: A Framework for Learning Robot-Agnostic Force-based Skills Involving Sustained Contact Object Manipulation*. In: *Proceedings of the 2025 IEEE International Conference on Robotics and Automation*, pp. 4782–4788, doi:10.1109/ICRA55743.2025.11127866.
- [4] Shangding Gu, Long Yang, Yali Du, Guang Chen, Florian Walter, Jun Wang & Alois Knoll (2024): *A Review of Safe Reinforcement Learning: Methods, Theories, and Applications*. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46(12), pp. 11216–11235, doi:10.1109/TPAMI.2024.3457538.
- [5] Neville Hogan (1984): *Impedance Control: An Approach to Manipulation*. In: *1984 American Control Conference*, pp. 304–313, doi:10.23919/ACC.1984.4788393.
- [6] O. Khatib (1987): *A unified approach for motion and force control of robot manipulators: The operational space formulation*. *IEEE Journal on Robotics and Automation* 3(1), pp. 43–53, doi:10.1109/JRA.1987.1087068.
- [7] Constantine Lignos, Vasumathi Raman, Cameron Finucane, Mitchell P. Marcus & Hadas Kress-Gazit (2015): *Provably correct reactive control from natural language*. *Auton. Robots* 38(1), pp. 89–105, doi:10.1007/S10514-014-9418-8.
- [8] Lars Lindemann & Dimos V. Dimarogonas (2019): *Control Barrier Functions for Signal Temporal Logic Tasks*. *IEEE Control Systems Letters* 3(1), pp. 96–101, doi:10.1109/LCSYS.2018.2853182.
- [9] Amartya Mukherjee, Ruikun Zhou, Haocheng Chang & Jun Liu (2024): *Harmonic Control Lyapunov Barrier Functions for Constrained Optimal Control with Reach-Avoid Specifications*, doi:10.48550/arXiv.2310.02869.
- [10] Muhammad Zakiyullah Romdlony & Bayu Jayawardhana (2016): *Stabilization with guaranteed safety using Control Lyapunov–Barrier Function*. *Automatica* 66, pp. 39–47, doi:10.1016/j.automatica.2015.12.011. Available at <https://www.sciencedirect.com/science/article/pii/S0005109815005439>.
- [11] Krishnan Srinivasan, Benjamin Eysenbach, Sehoon Ha, Jie Tan & Chelsea Finn (2020): *Learning to be Safe: Deep RL with a Safety Critic*, doi:10.48550/arXiv.2010.14603. Available at <http://arxiv.org/abs/2010.14603>.
- [12] Mohit Srinivasan & Samuel Coogan (2019): *Control Of Mobile Robots Using Barrier Functions Under Temporal Logic Specifications*. *CoRR* abs/1908.04903, doi:10.48550/arXiv.1908.04903. Available at <http://arxiv.org/abs/1908.04903>.

- [13] Wei Xiao, Calin A. Belta & Christos G. Cassandras (2021): *High Order Control Lyapunov-Barrier Functions for Temporal Logic Specifications*. In: *2021 American Control Conference (ACC)*, pp. 4886–4891, doi:10.23919/ACC50511.2021.9483028.
- [14] Hongchao Zhang, Zhizhen Qin, Sicun Gao & Andrew Clark (2024): *SEEV: Synthesis with Efficient Exact Verification for ReLU Neural Barrier Functions*, doi:10.48550/arXiv.2410.20326. Available at <https://arxiv.org/abs/2410.20326>.